

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/221299170>

A Comparative Study on Content-Based Music Genre Classification

Conference Paper · January 2003

DOI: 10.1145/860435.860487 · Source: DBLP

CITATIONS

304

READS

2,615

3 authors, including:



Mitsunori Ogiwara

University of Miami

262 PUBLICATIONS 7,417 CITATIONS

[SEE PROFILE](#)



Qi Li

Western Kentucky University

63 PUBLICATIONS 1,924 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



Scale invariant representation of imbalanced points [View project](#)



The complexity of boolean dynamical systems [View project](#)

A Comparative Study on Content-Based Music Genre Classification

Tao Li
Computer Science Dept.
University of Rochester
Rochester, NY 14627-0226
taoli@cs.rochester.edu

Mitsunori Ogihara
Computer Science Dept.
University of Rochester
Rochester, NY 14627-0226
ogihara@cs.rochester.edu

Qi Li
Dept. of CIS
University of Delaware
Newark, DE 19716
qili@cis.udel.edu

ABSTRACT

Content-based music genre classification is a fundamental component of music information retrieval systems and has been gaining importance and enjoying a growing amount of attention with the emergence of digital music on the Internet. Currently little work has been done on automatic music genre classification, and in addition, the reported classification accuracies are relatively low. This paper proposes a new feature extraction method for music genre classification, *DWCHs*¹. *DWCHs* capture the local and global information of music signals simultaneously by computing histograms on their Daubechies wavelet coefficients. Effectiveness of this new feature and of previously studied features are compared using various machine learning classification algorithms, including Support Vector Machines and Linear Discriminant Analysis. It is demonstrated that the use of *DWCHs* significantly improves the accuracy of music genre classification.

Categories and Subject Descriptors

H.3 [Information Storage and Retrieval]: Content Analysis and Indexing, Information Search and Retrieval; I.2 [Artificial Intelligence]: Learning; I.5 [Pattern Recognition]: Applications; J.5 [Arts and Humanities]: music

General Terms

Algorithms, Measurement, Performance, Experimentation, Verification

Keywords

Music Genre Classification, wavelet coefficients histogram, feature extraction, multi-class classification

¹*DWCHs* stands for Daubechies Wavelet Coefficient Histograms

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

SIGIR'03, July 28–August 1, 2003, Toronto, Canada.
Copyright 2003 ACM 1-58113-646-3/03/0007 ...\$5.00.

1. INTRODUCTION

Music is not only for entertainment and for pleasure, but has been used for a wide range of purposes due to its social and physiological effects. At the beginning of the 21st century the world is facing ever-increasing growth of the on-line music information, empowered by the permeation of Internet into daily life. Efficient and accurate automatic music information processing (accessing and retrieval, in particular) will be an extremely important issue, and it has been enjoying a growing amount of attention. Music can be classified based on its style and the styles have a hierarchical structure. A currently popular topic in automatic music information retrieval is the problem of organizing, categorizing, and describing music contents on the web. Such endeavor can be found in on-line music databases such as mp3.com and Napster. One important aspect of the the genre structures in these on-line databases is that the genre is specified by human experts as well as amateurs (such as the user) and that labeling process is time-consuming and expensive. Currently music genre classification is done mainly by hand because giving a precise definition of a music genre is extremely difficult and, in addition, many music sounds sit on boundaries between genres. These difficulties are due to the fact that music is an art that evolves, where performers and composers have been influenced by music in other genres. However, it has been observed that audio signals (digital or analog) of music belonging to the same genre share certain characteristics, because they are composed of similar types of instruments, having similar rhythmic patterns, and similar pitch distributions [7]. This suggests feasibility of automatic musical genre classification.

By automatic musical genre classification we mean here the most strict form of the problem, i.e., classification of music signals into a single unique class based computational analysis of music feature representations. Automatic music genre classification is a fundamental component of music information retrieval systems. We divide the process of genre categorization in music into two steps: *feature extraction* and *multi-class classification*. In the feature extraction step, we extract from the music signals information representing the music. The features extract should be *comprehensive* (representing the music very well), *compact* (requiring a small amount of storage), and *effective* (not requiring much computation for extraction). To meet the first requirement the design has to be made so that the both low-level and high-level information of the music is included. In the second step, we build a mechanism (an algorithm and/or

a mathematical model) for identifying the labels from the representation of the music sounds with respect to their features.

There has been a considerable amount of work in extracting features for speech recognition and music-speech discrimination, but much less work has been reported on the development of descriptive features specifically for music signals. To the best of our knowledge, currently the most influential approach to direct modeling of music signals for automatic genre classification is due to Tsanetakis and Cook [31], where the timbral texture, rhythm, and pitch content features are explicitly developed. The accuracy of classification based on these features, however, is only 61% is achieved on their ten-genre sound dataset. This raises the question of whether there are different features that are more useful in music classification and whether the use of statistical or machine learning techniques (e.g., discriminant analysis and support vector machines) can improve the accuracy. The main goal of this paper is to address these two issues.

In this paper, we propose a new feature extraction method based on wavelet coefficients histogram, which we call *DWCHs*. *DWCHs* represent both local and global information by computing histograms on Daubechies wavelet coefficients at different frequency subbands with different resolutions. Using *DWCHs* along with advanced machine learning techniques, accuracy of music genre classification has been significantly improved. On the dataset of [31], the classification has been increased to almost 80%. On genre specific classification, i.e., distinguish one genre from the rest, the accuracy can be as high as 98%. We also examined which features proposed in [31] are the most effective. It turns out that the timbral features combined with the Mel-frequency Cepstral Coefficients achieve a high accuracy on each of the multi-class classification algorithms we tested. The rest of the paper is organized as follows: Section 2 reviews the related work on automatic genre classification. Section 3 gives a brief overview of the previously proposed features. Section 4 presents the wavelet coefficients histogram features. Section 5 describes the machine learning techniques used for our experiments. Section 6 presents and discusses our experimental results. Finally, Section 7 provides our conclusions.

2. RELATED WORK

Many different features can be used for music classification, e.g., reference features including title and composer, content-based acoustic features including tonality, pitch, and beat, symbolic features extracted from the scores, and text-based features extracted from the song lyrics. In this paper we are interested in content-based features².

The content-based acoustic features are classified into timbral texture features, rhythmic content features, and pitch content features [31]. Timbral features are mostly originated from traditional speech recognition techniques. They are usually calculated for every short-time frame of sound based on the Short Time Fourier Transform (STFT) [24]. Typical timbral features include Spectral Centroid, Spectral Rolloff, Spectral Flux, Energy, Zero Crossings, Linear Prediction Coefficients, and Mel-Frequency Cepstral Coefficients (MFCCs) (see [24] for more detail). Among these tim-

bral features MFCCs have been dominantly used in speech recognition. Logan [18] examines MFCCs for music modeling and music/speech discrimination. Rhythmic content features contains information about the regularity of the rhythm, the beat and tempo information. Tempo and beat tracking from acoustic musical signals has been explored in [13, 15, 26]. Foote and Uchihashi [10] use the beat spectrum to represent rhythm. Pitch content features deals with the frequency information of the music bands and are obtained using various pitch detection techniques.

Much less work has been reported on music genre classification. Tzanetakis and Cook [31] proposes a comprehensive set of features for direct modeling of music signals and explores the use of those features for musical genre classification using K-Nearest Neighbors and Gaussian Mixture models. Lambrou et al. [14] uses statistical features in the temporal domain as well as three different wavelet transform domains to classify music into rock, piano and jazz. Deshpande et al. [5] uses Gaussian Mixtures, Support Vector Machines and Nearest Neighbors to classify the music into rock, piano, and jazz based on timbral features. Pye [23] investigates the use of Gaussian Mixture Modeling (GMM) and Tree-Based Vector Quantization in music genre classification. Soltau et al. [28] propose an approach of representing temporal structures of input signal. They show that this new set of abstract features can be learned via artificial neural networks and can be used for music genre identification.

The problem of discriminate music and speech has been investigated by Saunders [25], Scheier and Slaney [27]. Zhang and Kuo [34] propose a heuristic rule-based system to segment and classify audio signals from movies or TV programs. In [33] audio contents are divided into instrument sounds, speech sounds, and environment sounds using automatically extracted features. Foote [9] constructs a learning tree vector quantizer using twelve MFCCs plus energy as audio features for retrieval. Li and Khokhar [16] propose nearest feature line methods for content based classification audio retrieval.

3. THE CONTENT-BASED FEATURES

3.1 Timbral Textural Feature

Timbral textual features are used to differentiate mixture of sounds that are possibly with the same or similar rhythmic and pitch contents. The use of these features originates from speech recognition. To extract timbral features, the sound signals are first divided into frames that are statistically stationary, usually by applying a windowing function at fixed intervals. The window function, typically a Hamming window, removes edge effects. Timbral textual features are then computed for each frame and the statistical values (such as the mean and the variance) of those features are calculated.

- **Mel-Frequency Cepstral Coefficients (MFCCs)** are designed to capture short-term spectral-based features. After taking the logarithm of the amplitude spectrum based on STFT for each frame, the frequency bins are grouped and smoothed according to Mel-frequency scaling, which is design to agree with perception. MFCCs are generated by decorrelating the Mel-spectral vectors using discrete cosine transform.

²Our future work will address the issue of mixing different types of features.

- **Spectral Centroid** is the centroid of the magnitude spectrum of STFT and it is a measure of spectral brightness.
- **Spectral Rolloff** is the frequency below which 85% of the magnitude distribution is concentrated. It measures the spectral shape.
- **Spectral Flux** is the squared difference between the normalized magnitudes of successive spectral distributions. It measures the amount of local spectral change.
- **Zero Crossings** is the number of time domain zero crossings of the signal. It measures noisiness of the signal.
- **Low Energy** is the percentage of frames that have energy less than the average energy over the whole signal. It measures amplitude distribution of the signal.

3.2 Rhythmic Content Features

Rhythmic content features characterize the movement of music signals over time and contain such information as the regularity of the rhythm, the beat, the tempo, and the time signature. The feature set for representing rhythm structure is based on detecting the most salient periodicities of the signal and it is usually extracted from beat histogram. To construct the beat histogram, the time domain amplitude envelope of each band is first extracted by decomposing the music signal into a number of octave frequency band. Then, the envelopes of each band are summed together followed by the computation of the autocorrelation of resulting sum envelop. The dominant peaks of the autocorrelation function, corresponding to the various periodicities of the signal's envelope, are accumulated over the whole sound file into a beat histogram where each bin corresponds to the peak lag. The rhythmic content features are then extracted from the beat histogram, and generally they contain relative amplitude of the first and the second histogram peak, ration of the amplitude of the second peak divided by the amplitude of the first peak, periods of the first and second peak, overall sum of the histogram.

3.3 Pitch Content Features

The pitch content features describe the melody and harmony information about music signals and are extracted based on various pitch detection techniques. Basically the dominant peaks of the autocorrelation function, calculated via the summation of envelopes for each frequency band obtained by decomposing the signal, are accumulated into pitch histograms and the pitch content features are then extracted from the pitch histograms. The pitch content features typically include: the amplitudes and periods of maximum peaks in the histogram, pitch intervals between the two most prominent peaks, the overall sums of the histograms.

4. DWCHS OF MUSIC SIGNALS

It is not difficult to see that the traditional feature extraction described in Section 3 more or less capture incomplete information of music signals. Timbral textural features are standard features used in speech recognition and are calculated for every short-time frame of sound while rhythmic and pitch content features are computed over the whole file. In other words, timbral features capture the statistics of

local information of music signals from a global perspective, but not enough in representing the global information of the music. Moreover, as indicated by our experiments in Section 6, the rhythm and pitch content features don't seem to capture enough information content for classification purpose. In this section, we will propose a new feature extraction technique for music genre classification, namely, *DWCHs*, based on wavelet histogram to capture local and information of music signals simultaneously.

4.1 Wavelet Basics

The wavelet transform is a synthesis of ideas emerging over many years from different fields such as mathematics and image and signal processing and has been widely used in information retrieval and data mining [19, 20]. A complete survey on wavelet application in data mining can be found in [17]. Generally speaking, the wavelet transform, providing good time and frequency resolution, is a tool that divides up data, functions, or operators into different frequency components and then studies each component with a resolution matched to its scale [3]. Straightforwardly, a wavelet coefficients histogram is the histogram of the (rounded) wavelet coefficients obtained by convolving a wavelet filter with an input music signal (details on histogram and wavelet filter/analysis can be found in [3, 29] respectively).

Several favorable properties of wavelets, such as compact support, vanishing moments and decorrelated coefficients, make them useful tools for signal representation and transformation. Generally speaking, wavelets are designed to give good time resolution at high frequencies and good frequency resolution at low frequencies. Compact support guarantees the localization of wavelet, vanishing moment property allows wavelet focusing on most important information and discarding noisy signal, and decorrelated coefficients property enables wavelet to reduce temporal correlation so that the correlation of wavelet coefficients are much smaller than that of the corresponding temporal process [8]. Hence, after wavelet transform, the complex signal in the time domain can be reduced into a much simpler process in the wavelet domain. By computing the histograms of wavelet coefficients, we then could get a good estimation of the probability distribution over time. The good probability estimation thus leads to a good feature representation.

4.2 DWCHs

A sound file is a kind of oscillation waveform in time-domain and can be considered as a two dimensional entity of the amplitude over time, in the form of $M(t) = D(A, t)$ where A is the amplitude and generally ranges from $[-1, 1]$. The distinguishing characteristics are contained in the amplitude variation, and in consequence, identifying the amplitude variation would be essential for music categorization.

On one hand, the histogram technique is an efficient means for the distribution estimation. However, the raw signal in time domain is not good representation particularly for content-based categorization since the most distinguished characteristics are hidden in frequency domain. On the other hand, the sound frequency spectrum is generally divided into octaves with each having a unique quality. An octave is the interval between any frequencies that have a tonal ratio of 2 to 1, a logarithmic-relation in frequency band. The wavelet decomposition scheme matches

the models of sound octave division for perceptual scales and provides good time and frequency resolution [16]. In other words, the decomposition of audio signal using wavelets produces a set of subband signals at different frequencies corresponding different characteristics. This motivates the use of wavelet histogram techniques for feature extraction. The wavelet coefficients are distributed in various frequency bands at different resolutions.

There are many kinds of wavelet filters, including Daubechies wavelet filter, Gabor filter etc. Daubechies wavelet filters are the one commonly in image retrieval (more details on wavelet filter can be found in [3]). In our work, we use Daubechies wavelet filter Db8 with seven levels of decomposition. After the decomposition, we construct the histogram of the wavelet coefficients at each subband. The coefficient histogram provides a good approximation of the waveform variations at each subband. From probability theory, a probability distribution is uniquely characterized by its moments. Hence, if we interpret the waveform distribution as a probability distribution, then it can be characterized by its moments. To characterize the waveform distribution, the first three moments of a histogram is used [4]. The first three moments are the average, the variance and the skewness of each subband. We also compute the subband energy, defined as the mean of the absolute value of coefficients, for each subband. In addition, our final *DWCHs* feature set includes the tradition timbral features used for speech recognition.

Each music file in our datasets is 30-second signal that will be first converted to an extremely long vector. Based on an intuition of “self-similarity” of a music, i.e., its repeated theme, The *DWCHs* feature can be extracted on a small slice of an input music signal and in our experiments 3 seconds sound clip are used. In summary, our algorithm of *DWCHs* extraction contains the following steps:

1. The wavelet decomposition of the music signals is obtained
2. Histogram of each subband is constructed
3. Compute the first three moments of all histograms
4. Compute the subband energy for each subband

The algorithm is very easy to implement in Matlab that contains a complete wavelet package.

4.3 Examples

Concrete examples will make clearer the *DWCHs* features and their putative advantages. In Figure 1, from left to right and top to bottom, *DWCHs* features of ten music sounds drawn from the different genres of the ten-genre dataset (the Dataset A that we used in our experiments, to be described in Section 6) are shown. We observe the distinct difference of the feature representations of different music genres. For example, the features on several subbands, corresponding the features from 20 to 30, are clearly different for country and disco. The differences are quite consistent with the human feeling on music. Figure 2 and Figure 3 show *DWCHs* features of ten Blues and ten Ten classical music sounds taken from Dataset A respectively. We observe that similar features are present inside a single music genre. If our observations are correct, a unique *DWCHs* feature pattern exists in each music genre, and we can hypothesize that the use of *DWCHs* will improve classification of music genre.

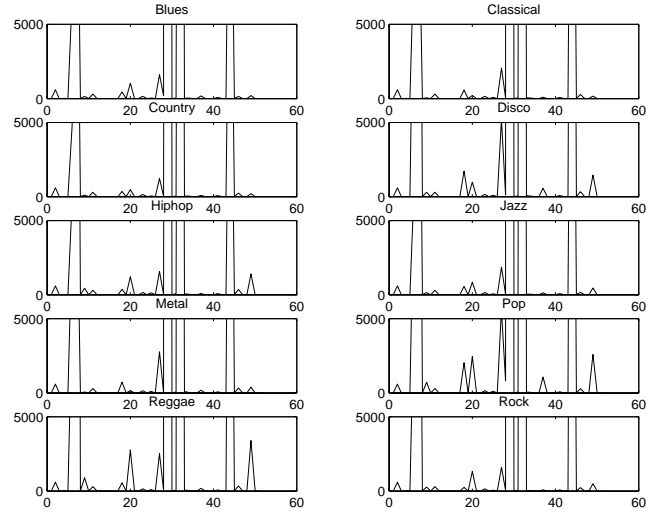


Figure 1: *DWCHs* of 10 music signals in different genre. The feature representations of different genre are mostly different to each other.

5. MULTI-CLASS LEARNING METHODS

Once the feature set has been extracted, the music genre classification problem is reduced to a multi-class classification problem. The problem can be formally defined as follows: The input to the problem is a set of training samples of the form of $\langle x_i, l_i \rangle$, where x_i is a data point and l_i is its label, chosen from a finite set of labels $\{c_1, c_2, \dots, c_K\}$. In our case the labels are music genres. The goal is to infer a function f that well approximates the mapping the x 's to their labels. Generally speaking, approaches to multi-class classification problems can be roughly divided into two groups. The first group consists of those binary classification algorithms that can be naturally extended to handle multi-class cases. This group contains such algorithm as Discriminant Analysis, K-Nearest Neighbors, regression, and decision trees including C4.5 and CART. The second group consists of methods that involve decomposition of multi-class classification problems to a collection binary ones. Many decomposition technique exists, including such popular methods as the one-versus-the-rest method, pairwise comparison, Error-Correcting Output Coding (ECOC), and multi-class objective functions.

The idea of the one-versus-the-rest method is as follows: to get a K -class classifier, first construct a set of binary classifiers C_1, C_2, \dots, C_K . Each binary classifier is first trained to separate one class from the rest and then the multi-class classification is carried out according to the maximal output of the binary classifiers. In pairwise comparison, a classifier is trained for each possible pair of classes. For K classes, this results in $(K-1)K/2$ binary classifiers. Given a new instance, the multi-class classification is then executed by evaluating all $(K-1)K/2$ individual classifiers and assigning the instance to the class which gets the highest number of votes. Error-Correcting Output Coding (ECOC) [6], roughly speaking, generates a number of binary classification problems by intelligently splitting the original set of classes into two sets. In other words, each class is assigned a unique binary strings of length l (these strings are regarded

to codewords). Then l classifiers are trained to predict each bit of the string. For new instances, the predicted class is the one whose codeword is the closest (in Hamming distance) to the codeword produced by the classifiers. One-versus-the-rest and pairwise comparison can be regarded as two special cases of ECOC with specific coding schemes. Multi-class objective functions aims to directly modify the objective function of binary SVMs in such a way that it simultaneously allows the computation of a multi-class classifier. In practice, the choice of reduction method from multi-class to binary is problem-dependent and not a trivial task and it is fair to say that there is probably no method generally outperforms the others [1].

We conducted experiments with the following additional multi-class classification approaches (see [21] for more information about the methods):

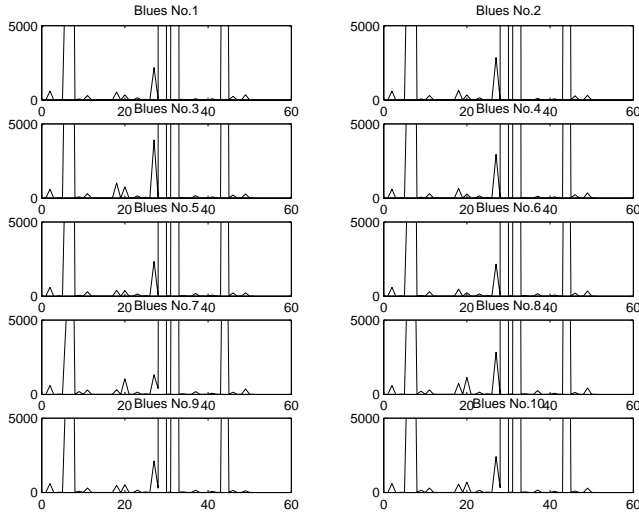


Figure 2: DWCHs of ten blues songs. The feature representations are similar.

- **Support Vector Machines** Support vector machines (SVMs) [32] have shown superb performance in binary classification tasks. Basically, Support Vector Machine aim at searching for a hyperplane that separates the positive data points and the negative data points with maximum margin. To extend SVMs for multi-class classification, we use one-versus-the-rest, pairwise comparison, and multi-class objective functions.
- **K-Nearest Neighbor (KNN)** KNN is a nonparametric classifier. It is proven that the error of KNN is asymptotically at most twice as large as the Bayesian error rate. KNN has been applied in various musical analysis problems. The basic idea is to allow a small number of neighbors to influence the decision on a point.
- **Gaussian Mixture Models (GMM)** GMM models the distribution of acoustics and has been widely used in music information retrieval. For each class we assume the existence of a probability density function (pdf) expressible as a mixture of a number of multi-dimensional Gaussian distributions. The iterative EM

algorithm is then used to estimate the parameters for each Gaussian component and the mixture weight.

- **Linear Discriminant Analysis (LDA)** In the statistical pattern recognition literature discriminant analysis approaches are known to learn discriminative feature transformations very well. The approach has been successfully used in many classification tasks [11]. The basic idea of LDA is to find a linear transformation that best discriminates among classes and perform classification in the transformed space based on some metric such as Euclidean distances. Fisher discriminant analysis finds discriminative feature transform as eigenvectors of matrix $T = \hat{\Sigma}_w^{-1} \hat{\Sigma}_b$ where $\hat{\Sigma}_w$ is the intra-class covariance matrix and $\hat{\Sigma}_b$ is the inter-class covariance matrix. This matrix T captures both compactness of each class and separations between classes. So, the eigenvectors corresponding to the largest eigenvalues of T are expected to constitute a discriminative feature transform.

6. EXPERIMENTAL RESULTS

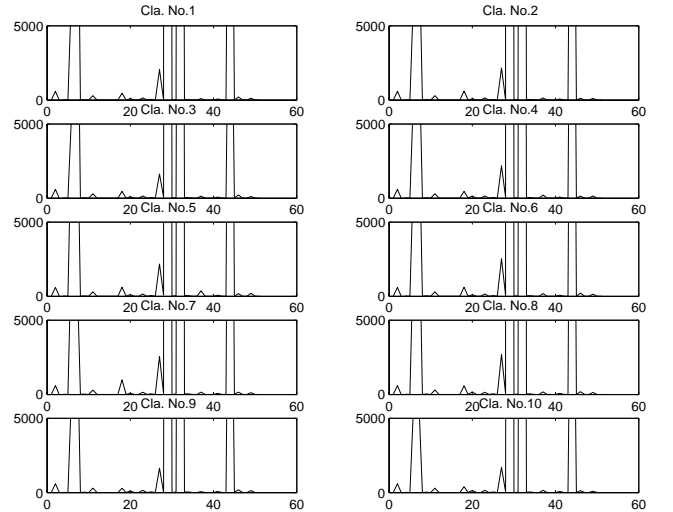


Figure 3: DWCHs of ten classical songs.

6.1 Datasets

We used two datasets for our experiments. The first dataset, Dataset A, contains 1000 songs over ten genres with 100 songs per genre. This dataset was used in [31]. The ten genres are Blues, Classical, Country, Disco, Hiphop, Jazz, Metal, Pop, Reggae, and Rock. The excerpts of the dataset were taken from radio, compact disks, and MP3 compressed audio files. The second dataset, Dataset B, contains 756 sounds over five genres: Ambient, Classical, Fusion, Jazz, and Rock. This dataset was constructed for this paper from the CD collection of the second author. The collection of 756 sound files was created from 189 music albums as follows: From each album the first four music tracks were chosen (three tracks from albums with only three music tracks). Then from each music track the sound signals over a period of 30 seconds after the initial 30 seconds were extracted in MP3. The distribution of different genres is: Ambient (109

files), Classical (164 files), Fusion (136 files), Jazz (251 files) and Rock (96 files)³. For both datasets, the sound files are converted to 22050Hz, 16-bit, mono audio files.

6.2 Experiment Setup

We used MARSYAS⁴, a public software framework for computer audition applications, for extracting the features proposed in [31]: *MFCCs*, *FFT*, *Beat* and *Pitch*. Mel-frequency Cepstral Coefficients (denoted by MFCCs), the timbral texture features excluding MFCCs (denoted by FFT), the rhythm content features (denoted by Beat), and the pitch contents feature (denoted by Pitch) (see Section 3). The *MFCCs* feature vector consists of the mean and variance of each of the first five MFCC coefficients over the frames, the *FFT* feature vector consists of the mean and variance of spectral centroid, of rolloff, of flux, and of zero crossings, and of low energy, the *Beat* feature vector consists of six features from the rhythm histogram, the *Pitch* feature vector consists of five features from pitch histograms. More information of the feature extraction can be found in [30]. Our original *DWCH* feature set contains four features for each of seven frequency subbands along with nineteen traditional timbral features. However, we found that not all the frequency subbands are informative and we only use selective subbands, resulting a feature vector of length 35.

For classification methods, we use three different reduction methods to extend SVM for multi-class: pairwise, one-against-the-rest, and multi-class objective functions. For one-against-the-rest and pairwise methods, our SVM implementation was based on the LIBSVM [2], a library for support vector classification and regression. For multi-class objective functions, our implementation was based on MPSVM [12]. For experiments involving SVMs, we tested them with linear, polynomial, and radius-based kernels and the best results are reported in the tables. For Gaussian Mixture Models, we used three Gaussian mixtures to model each music genre. For K-Nearest Neighbors, we set $k = 5$.

6.3 Results and Analysis

Table 1 shows the accuracy of the various classification algorithms on Dataset A. The bottom four rows show how the classifiers performed on a single set of features proposed in [31]. The experiments verify the fact that each of the traditional features contains useful yet incomplete information characterizing music signals. The classification accuracy on any single feature set is significantly better than random guess (the accuracy of random guess on dataset A is 10%). The performance with either FFT or MFCC was significantly higher than that with Beat or Pitch in each of the methods tested. This naturally raises a question of whether FFT and MFCC are each more suitable than Beat or Pitch for music genre classification. We combined the four sets of features in every possible way to examine the accuracy. The accuracy with only Beat and Pitch is significantly smaller than the accuracy with any combination that includes either FFT or MFCC. Indeed, the accuracy with only FFT and MFCC is almost the same as that with all four for all methods. This seems to answer positively to our question.

³A list of music sources for Dataset B is available at <http://www.cs.rochester.edu/u/ogihara/music/SIGIR-list.xls>.

⁴MARSYAS can be downloaded from <http://www.cs.princeton.edu/~gtzan/marsyas.html>.

On the other hand, the use of *DWCHs* further improved the accuracy on all methods. In particular, there is a significant jump in the accuracy when Support Vector Machines are used with either the pairwise or the one-versus-all approach. The accuracy of the one-versus-the-rest SVM is 78.5% on the average in the ten-fold cross validation. For some of the cross validation tests the accuracy went beyond 80%. This is a remarkable improvement over Tsanetakis and Cook's 61%. Perrot and Gjerdigen [22] report a human study in which college students were trained to learn a music company's genre classification on a ten-genre data collection in which about 70% of accuracy is achieved. Although these results are not directly comparable due to the different dataset collections, it clearly implies that the automatic content-based genre classification could possibly achieve similar accuracy as human performance. In fact, in comparison the performance of our best method seems to go far beyond that.

There are papers reporting better accuracy of automatic music genre recognition of smaller datasets. Pye [23] reports 90% on a total set of 175 songs covering six genres (Blues, Easy Listening, Classical, Opera, Dance, and Indie Rock). Soltau et al. [28] report 80% accuracy on four classes (Rock, Pop, Techno, and Classical). Just for the sake of comparison, we show in Table 2 the performance of the best classifier (*DWCHs* with SVM) on one-versus-all tests on each of the ten music genres in Dataset A. The performance of these classifiers are extremely good. Also, in Table 3 we show the performance of the multi-class classification for distinction among smaller numbers of classes. The accuracy gradually decreases as the number of classes increases.

Table 4 presents the results on our own dataset. This dataset was generated with little control by blindly taking 30 seconds after introductory 30 seconds of each music and covers many different albums, so the performance was anticipated to be lower than that for Database A. Also, there is the genre of Ambient, which covers music bridging between Classical and Jazz. The difficulty in classifying such borderline cases is compensated for the reduction in the number of classes. The overall performance was only 4 to 5% lower than that for Database A.

We observe that SVMs are always the best classifiers for content-based music genre classification. However, the choice of the reduction method from multi-class to binary seems to be problem-dependent and there is no clear overall winner. It is fair to say that there is probably no reduction method generally outperforms the others. Feature extraction is crucial for music genre classification. The choice of features is more important than the choice of classifiers. The variations of classification accuracy on different classification techniques are much smaller than those of different feature extraction techniques.

7. CONCLUSIONS AND FUTURE WORK

In this paper we proposed *DWCHs*, a new feature extraction method for music genre classification. *DWCHs* represent music signals by computing histograms on Daubechies wavelet coefficients at various frequency bands and it has significantly improved the classification accuracy. We also provide a comparative study of various feature extraction and classification methods and investigate the classification performance of various classification methods on different feature sets. To the best of our knowledge, this is the first

Features	Methods					
	SVM1	SVM2	MPSVM	GMM	LDA	KNN
<i>DWCHs</i>	74.9(4.97)	78.5(4.07)	68.3(4.34)	63.5(4.72)	71.3(6.10)	62.1(4.54)
Beat+FFT+MFCC+Pitch	70.8(5.39)	71.9(5.09)	66.2(5.23)	61.4(3.87)	69.4(6.93)	61.3(4.85)
Beat+FFT+MFCC	71.2(4.98)	72.1(4.68)	64.6(4.16)	60.8(3.25)	70.2(6.61)	62.3(4.03)
Beat+FFT+Pitch	65.1(4.27)	67.2(3.79)	56.0(4.67)	53.3(3.82)	61.1(6.53)	51.8(2.94)
Beat+MFCC+Pitch	64.3(4.24)	63.7(4.27)	57.8(3.82)	50.4(2.22)	61.7(5.23)	54.0(3.30)
FFT+MFCC+Pitch	70.9(6.22)	72.2(3.90)	64.9(5.06)	59.6(3.22)	69.9(6.76)	61.0(5.40)
Beat+FFT	61.7(5.12)	62.6(4.83)	50.8(5.16)	48.3(3.82)	56.0(6.73)	48.8(5.07)
Beat+MFCC	60.4(3.19)	60.2(4.84)	53.5(4.45)	47.7(2.24)	59.6(4.03)	50.5(4.53)
Beat+Pitch	42.7(5.37)	41.1(4.68)	35.6(4.27)	34.0(2.69)	36.9(4.38)	35.7(3.59)
FFT+MFCC	70.5(5.98)	71.8(4.83)	63.6(4.71)	59.1(3.20)	66.8(6.77)	61.2(7.12)
FFT+Pitch	64.0(5.16)	68.2(3.79)	55.1(5.82)	53.7(3.15)	60.0(6.68)	53.8(4.73)
MFCC+Pitch	60.6(4.54)	64.4(4.37)	53.3(2.95)	48.2(2.71)	59.4(4.50)	54.7(3.50)
Beat	26.5(3.30)	21.5(2.71)	22.1(3.04)	22.1(1.91)	24.9(2.99)	22.8(5.12)
FFT	61.2(6.74)	61.8(3.39)	50.6(5.76)	47.9(4.91)	56.5(6.90)	52.6(3.81)
MFCC	58.4(3.31)	58.1(4.72)	49.4(2.27)	46.4(3.09)	55.5(3.57)	53.7(4.11)
Pitch	36.6(2.95)	33.6(3.23)	29.9(3.76)	25.8(3.02)	30.7(2.79)	33.3(3.20)

Table 1: Classification accuracy of the learning methods tested on Dataset A using various combinations of features. The accuracy values are calculated via ten-fold cross validation. The numbers within parentheses are standard deviations. SVM1 and SVM2 respectively denote the pairwise SVM and the one-versus-the-rest SVM.

Number	Genre	Accuracy	4	Disco	92.69 (2.54)	8	Hiphop	96.49 (1.28)
1	Blues	95.49 (1.27)	5	Jazz	97.90(0.99)	9	Reggae	92.30 (2.49)
2	Classical	98.89 (1.10)	6	Metal	95.29 (2.18)	10	Rock	91.29 (2.96)
3	Country	94.29 (2.49)	7	Pop	95.80 (1.69)			

Table 2: Genre specific accuracy of SVM1 on DWCHs. The results are calculated via ten fold cross validation and each entry in the table is in the form of accuracy(standard deviation).

Classes	Methods					
	SVM1	SVM2	MPSVM	GMM	LDA	KNN
1 & 2	98.00(3.50)	98.00(2.58)	99.00(2.11)	98.00(3.22)	99.00(2.11)	97.5(2.64)
1, 2 & 3	92.33(5.46)	92.67(4.92)	93.33(3.51)	91.33(3.91)	94.00(4.10)	87.00(5.54)
1 through 4	90.5(4.53)	90.00(4.25)	89.75(3.99)	85.25(5.20)	89.25(3.92)	83.75(5.92)
1 through 5	88.00(3.89)	86.80(4.54)	83.40(5.42)	81.2(4.92)	86.2(5.03)	78.00(5.89)
1 through 6	84.83(4.81)	86.67(5.27)	81.0(6.05)	73.83(5.78)	82.83(6.37)	73.5(6.01)
1 through 7	83.86(4.26)	84.43(3.53)	78.85(3.67)	74.29(6.90)	81.00(5.87)	73.29(5.88)
1 through 8	81.5(4.56)	83.00(3.64)	75.13(4.84)	72.38(6.22)	79.13(6.07)	69.38(5.47)
1 through 9	78.11(4.83)	79.78(2.76)	70.55(4.30)	68.22(7.26)	74.47(6.22)	65.56(4.66)

Table 3: Accuracy on various subsets of Dataset A using *DWCHs*. The class numbers correspond to those of Table 2. The accuracy values are calculated via ten-fold cross validation. The numbers in the parentheses are the standard deviations.

Features	Methods					
	SVM1	SVM2	MPSVM	GMM	LDA	KNN
<i>DWCHs</i>	71.48(6.84)	74.21(4.65)	67.16(5.60)	64.77(6.15)	65.74(6.03)	61.84(4.88)
Beat+FFT+MFCC+Pitch	68.65(3.90)	69.19(4.32)	65.21(3.63)	63.08(5.89)	66.00(5.57)	60.59(5.43)
FFT+MFCC	66.67(4.40)	70.63(4.13)	64.29(4.54)	61.24(6.29)	65.35(4.86)	60.78(4.30)
Beat	43.37(3.88)	44.52(4.14)	41.01(4.46)	37.95(5.21)	40.87(4.50)	41.27(2.96)
FFT	61.65(5.57)	62.19(5.26)	54.76(2.94)	50.80(4.89)	57.94(5.11)	57.42(5.64)
MFCC	60.45(5.12)	67.46(3.57)	57.42(4.67)	53.43(5.64)	59.26(4.77)	59.93(3.49)
Pitch	37.56(4.63)	39.37(3.88)	36.49(5.12)	29.62(5.89)	37.82(4.67)	38.89(5.04)

Table 4: Classification accuracy of the learning methods tested on Dataset B using various combinations of features calculated via via ten-fold cross validation. The numbers within parentheses are standard deviations.

study of such kind in music genre classification.

Acknowledgments

The authors thank George Tzanetakis for useful discussions and for kindly sharing his data with us. This work is supported in part by NSF grants EIA-0080124, DUE-9980943, and EIA-0205061, and in part by NIH grants RO1-AG18231 (5-25589) and P30-AG18254.

8. REFERENCES

- [1] E. L. Allwein, R. E. Schapire, and Y. Singer. Reducing multiclass to binary: A unifying approach for margin classifiers. In *Proc. 17th International Conf. on Machine Learning*, pages 9–16. Morgan Kaufmann, San Francisco, CA, 2000.
- [2] C.-C. Chang and C.-J. Lin. *LIBSVM: a library for support vector machines*, 2001. Software available at <http://www.csie.ntu.edu.tw/~cjlin/libsvm>.
- [3] I. Daubechies. *Ten lectures on wavelets*. SIAM, Philadelphia, 1992.
- [4] A. David and S. Panchanathan. Wavelet-histogram method for face recognition. *Journal of Electronic Imaging*, 9(2):217–225, 2000.
- [5] H. Deshpande, R. Singh, and U. Nam. Classification of music signals in the visual domain. In *Proceedings of the COST-G6 Conference on Digital Audio Effects*, 2001.
- [6] T. G. Dietterich and G. Bakiri. Solving multiclass learning problems via error-correcting output codes. *Journal of Artificial Intelligence Research*, 2:263–286, 1995.
- [7] W. J. Dowling and D. L. Harwood. *Music Cognition*. Academic Press, Inc, 1986.
- [8] P. Flandrin. Wavelet analysis and synthesis of fractional Brownian motion. *IEEE Transactions on Information Theory*, 38(2):910–917, 1992.
- [9] J. Foote. Content-based retrieval of music and audio. In *In Multimedia Storage and Archiving Systems II, Proceedings of SPIE*, pages 138–147, 1997.
- [10] J. Foote and S. Uchihashi. The beat spectrum: a new approach to rhythm analysis. In *IEEE International Conference on Multimedia & Expo 2001*, 2001.
- [11] K. Fukunaga. *Introduction to statistical pattern recognition*. Academic Press, New York, 2nd edition, 1990.
- [12] G. Fung and O. L. Mangasarian. Multicategory proximal support vector machine classifiers. Technical Report 01-06, University of Wisconsin at Madison, 2001.
- [13] M. Goto and Y. Muraoka. A beat tracking system for acoustic signals of music. In *ACM Multimedia*, pages 365–372, 1994.
- [14] T. Lambrou, P. Kudumakis, R. Speller, M. Sandler, and A. Linney. Classification of audio signals using statistical features on time and wavelet transform domains. In *Proc. Int. Conf. Acoustic, Speech, and Signal Processing (ICASSP-98)*, volume 6, pages 3621–3624, 1998.
- [15] J. Laroche. Estimating tempo, swing and beat locations in audio recordings. In *Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA01)*, 2001.
- [16] G. Li and A. A. Khokhar. Content-based indexing and retrieval of audio data using wavelets. In *IEEE International Conference on Multimedia and Expo (II)*, pages 885–888, 2000.
- [17] T. Li, Q. Li, S. Zhu, and M. Ogihara. A survey on wavelet applications in data mining. *SIGKDD Explorations*, 4(2):49–68, 2003.
- [18] B. Logan. Mel frequency cepstral coefficients for music modeling. In *Proc. Int. Symposium on Music Information Retrieval/ISMIR*, 2000.
- [19] M. K. Mandal, T. Aboulnasr, and S. Panchanathan. Fast wavelet histogram techniques for image indexing. *Computer Vision and Image Understanding: CVIU*, 75(1–2):99–110, 1999.
- [20] Y. Matias, J. S. Vitter, and M. Wang. Wavelet-based histograms for selectivity estimation. In *Proceeding of the ACM SIGMOD Conference*, pages 448–459, 1998.
- [21] T. M. Mitchell. *Machine Learning*. The McGraw-Hill Companies, Inc., 1997.
- [22] D. Perrot and R. R. Gjerdingen. Scanning the dial: an exploration of factors in the identification of musical style. In *Proceedings of the 1999 Society for Music Perception and Cognition*, page 88, 1999.
- [23] D. Pye. Content-based methods for managing electronic music. In *Proceedings of the 2000 IEEE International Conference on Acoustic Speech and Signal Processing*, 2000.
- [24] L. Rabiner and B. Juang. *Fundamentals of Speech Recognition*. Prentice-Hall, NJ, 1993.
- [25] J. Saunders. Real-time discrimination of broadcast speech/music. In *Proc. ICASSP 96*, pages 993–996, 1996.
- [26] E. Scheirer. Tempo and beat analysis of acoustic musical signals. *Journal of the Acoustical Society of America*, 103(1), 1998.
- [27] E. Scheirer and M. Slaney. Construction and evaluation of a robust multifeature speech/music discriminator. In *Proc. ICASSP '97*, pages 1331–1334, Munich, Germany, 1997.
- [28] H. Soltau, T. Schultz, and M. Westphal. Recognition of music types. In *Proceedings of the 1998 IEEE International Conference on Acoustics, Speech and Signal Processing*, 1998.
- [29] M. Swain and D. Ballard. Color indexing. *Int. J. computer vision*, 7:11–32, 1991.
- [30] G. Tzanetakis and P. Cook. Marsyas: A framework for audio analysis. *Organized Sound*, 4(3):169–175, 2000.
- [31] G. Tzanetakis and P. Cook. Musical genre classification of audio signals. *IEEE Transactions on Speech and Audio Processing*, 10(5), July 2002.
- [32] V. N. Vapnik. *Statistical Learning Theory*. Wiley, New York, 1998.
- [33] E. Wold, T. Blum, D. Keislar, and J. Wheaton. Content-based classification, search and retrieval of audio. *IEEE Multimedia*, 3(2):27–36, 1996.
- [34] T. Zhang and C.-C. J. Kuo. Audio content analysis for online audiovisual data segmentation and classification. *IEEE Transactions on Speech and Audio Processing*, 3(4), 2001.